

# Router-based stream processing implemented in P4

Sammy Moseley

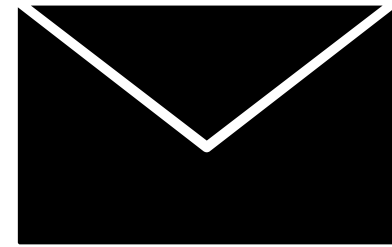
# Outline

1. Background on stream processing
2. Implementation using P4
3. Next steps

# Stream Processing Applications



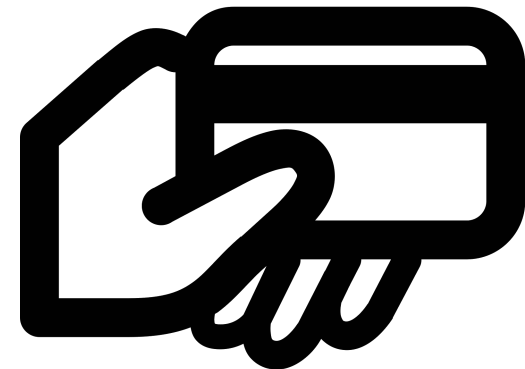
**Friend/Follow Requests**



**Messaging**

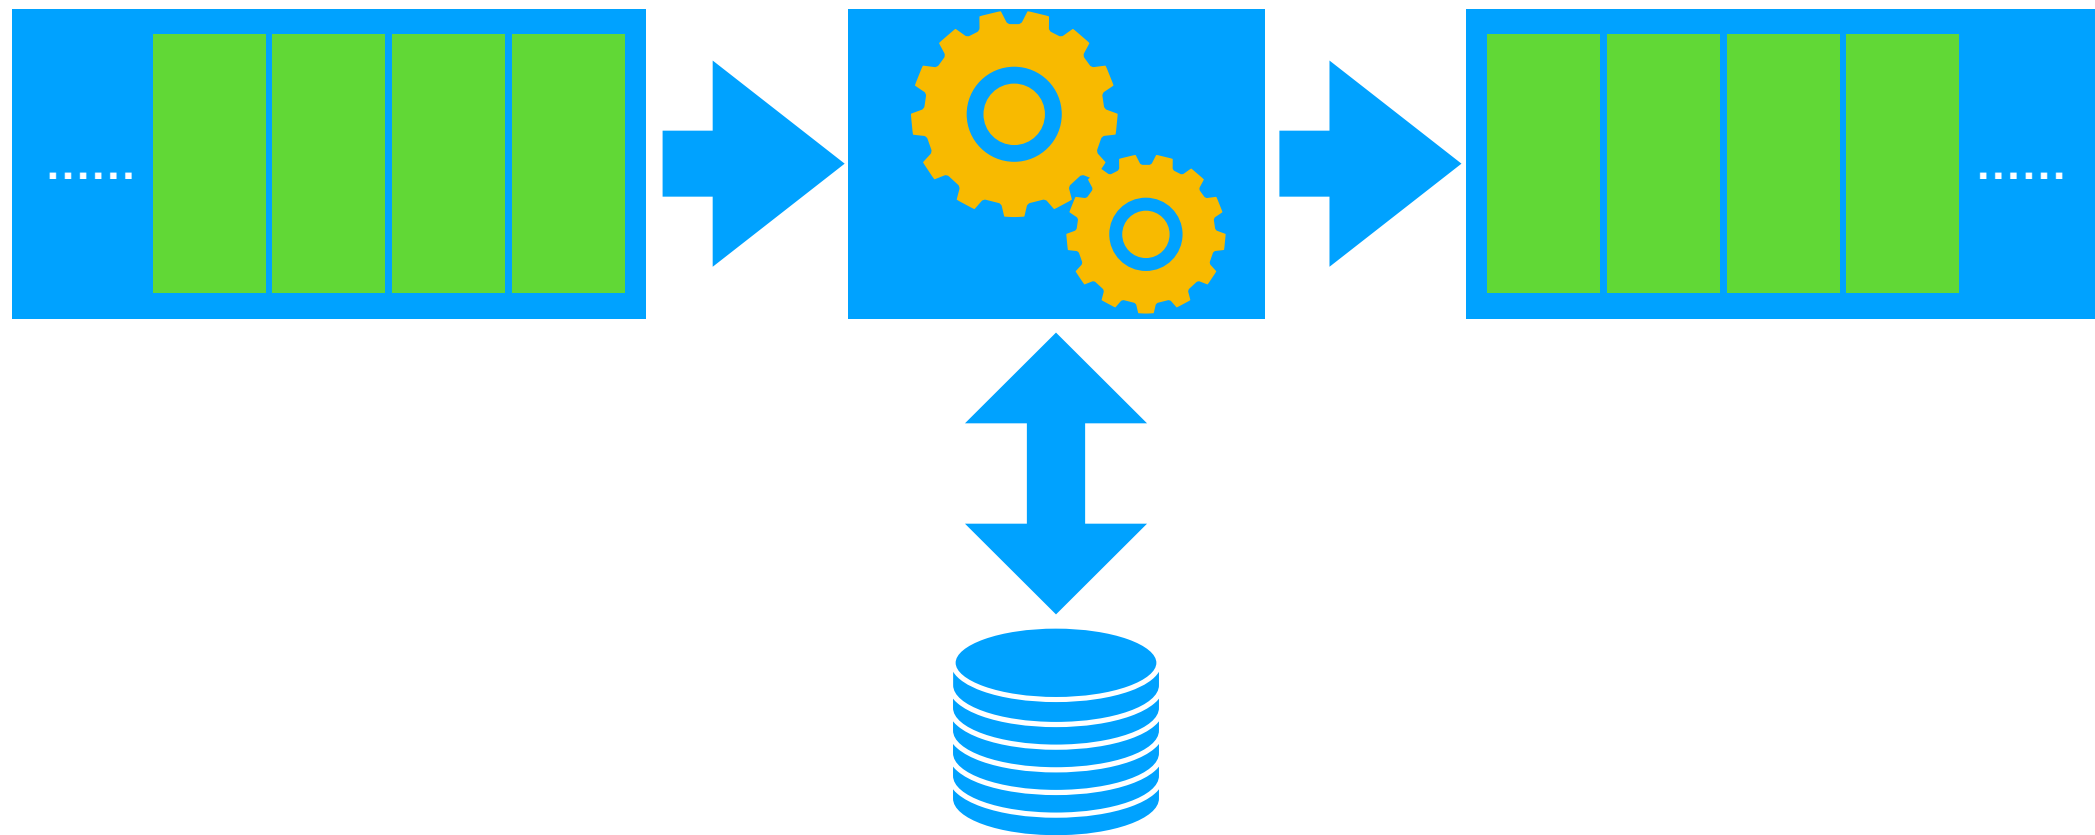


**Webpage view/post clicks**

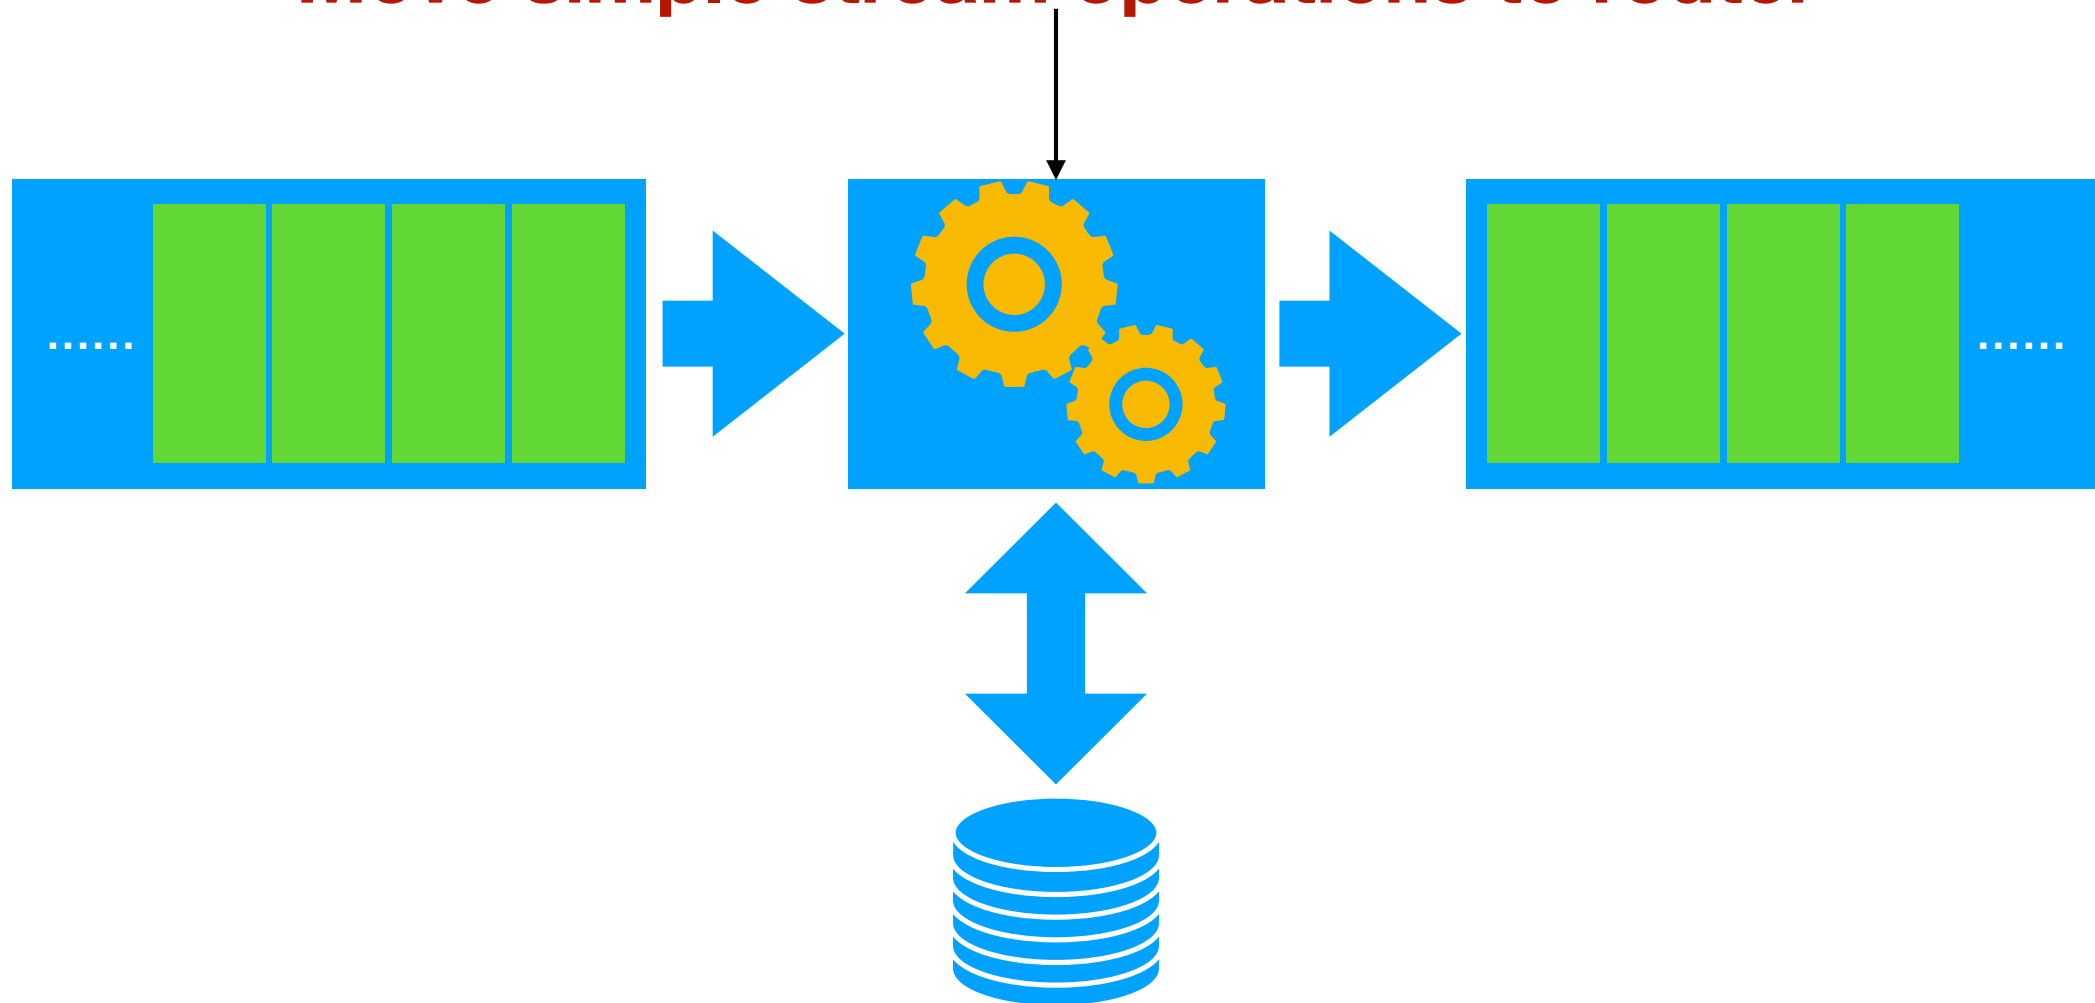


**Payment Transactions**

# A Quick Background on Stream Processing



**Move simple stream operations to router**



# Fundamental stream query operators

1. Filter
2. Map
3. Join
4. Aggregate
5. Zip

# Fundamental stream query operators

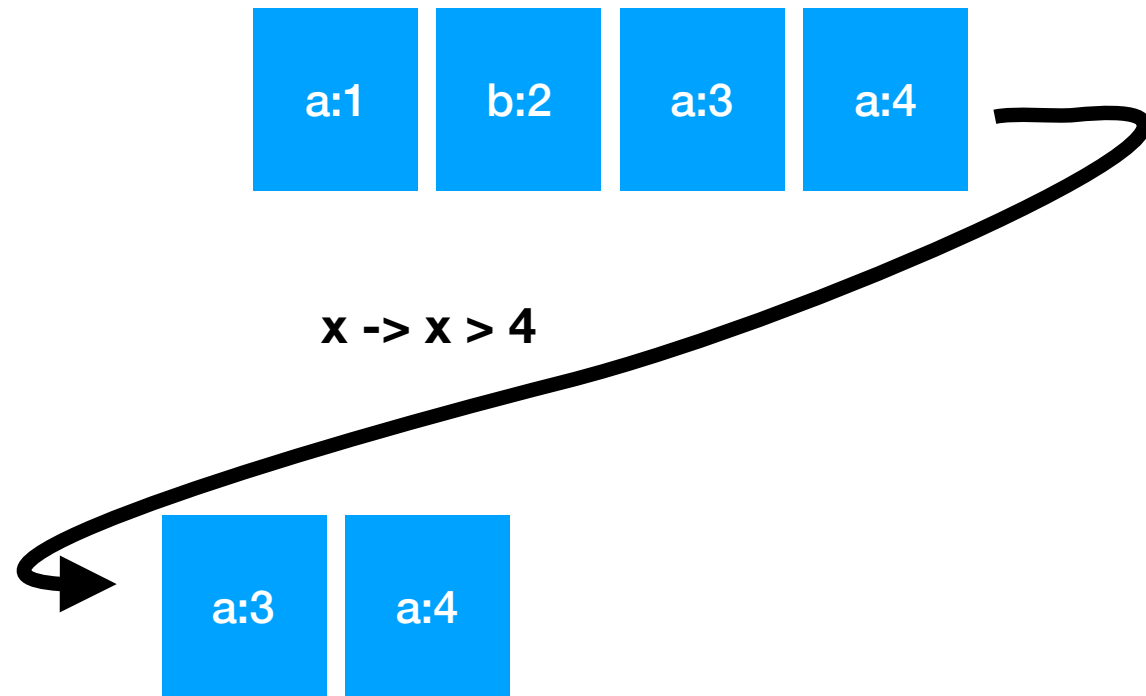
1. **Filter**

2. Map

3. Join

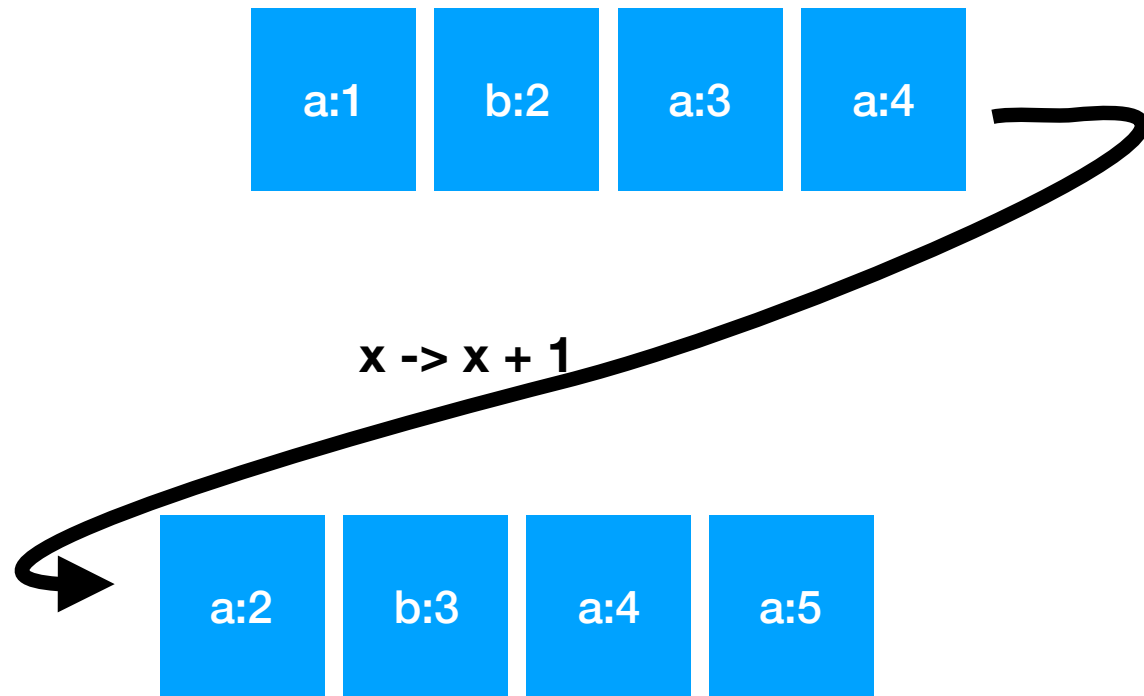
4. Aggregate

5. Zip



# Fundamental stream query operators

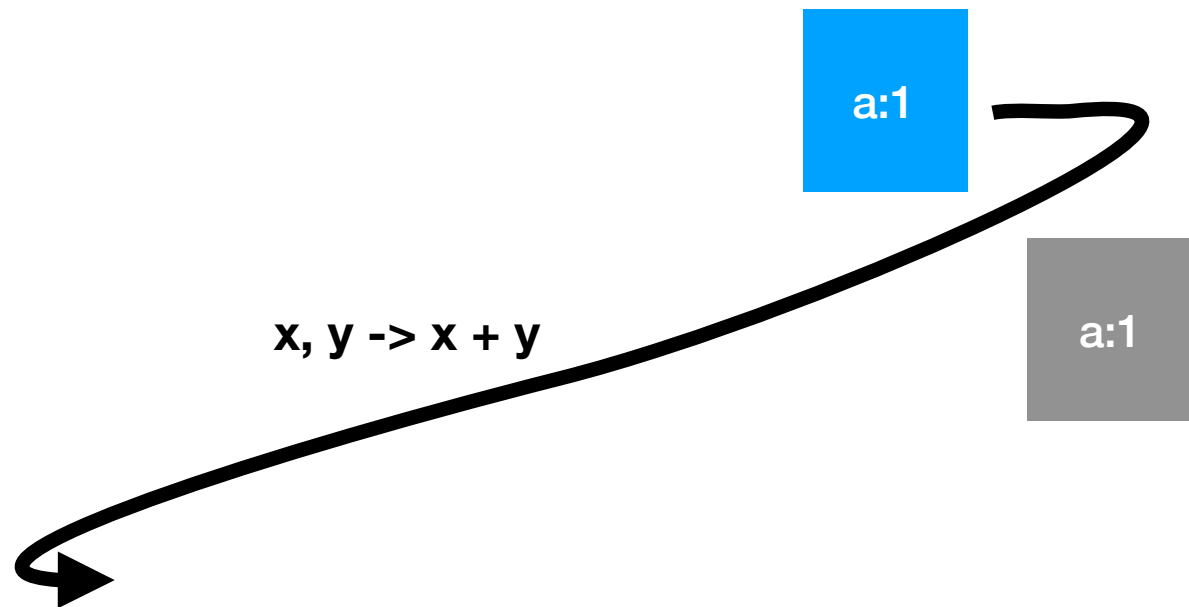
1. Filter
- 2. Map**
3. Join
4. Aggregate
5. Zip





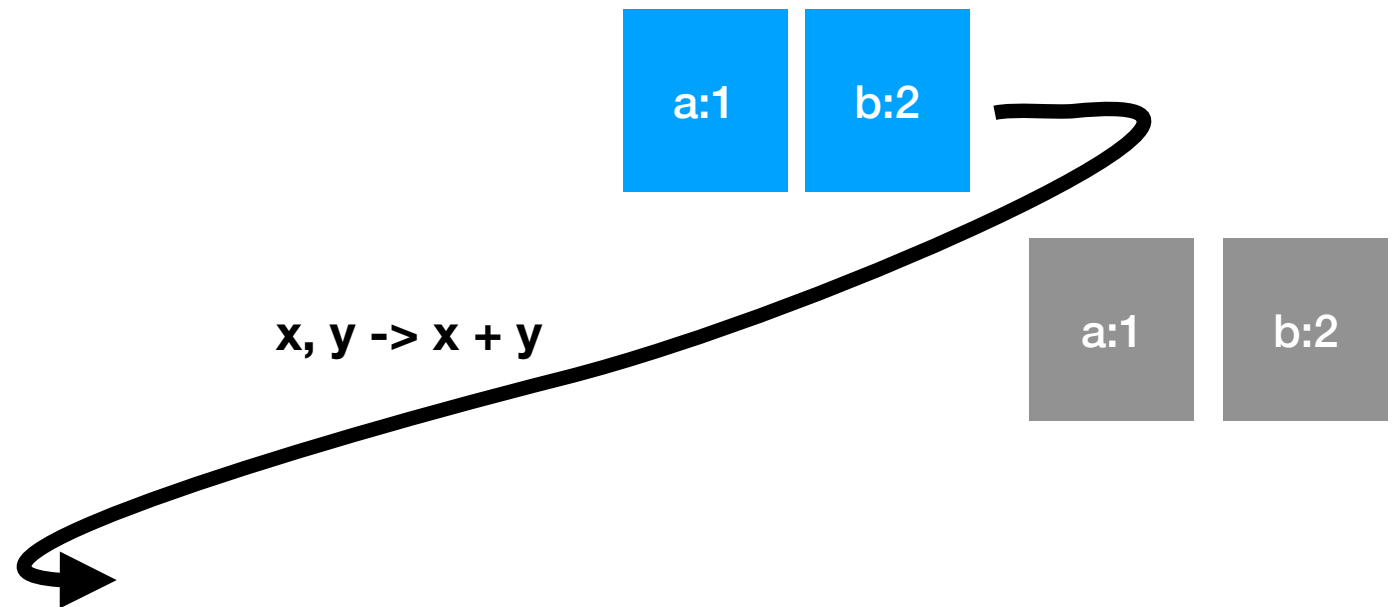
# Fundamental stream query operators

1. Filter
2. Map
- 3. Join**
4. Aggregate
5. Zip



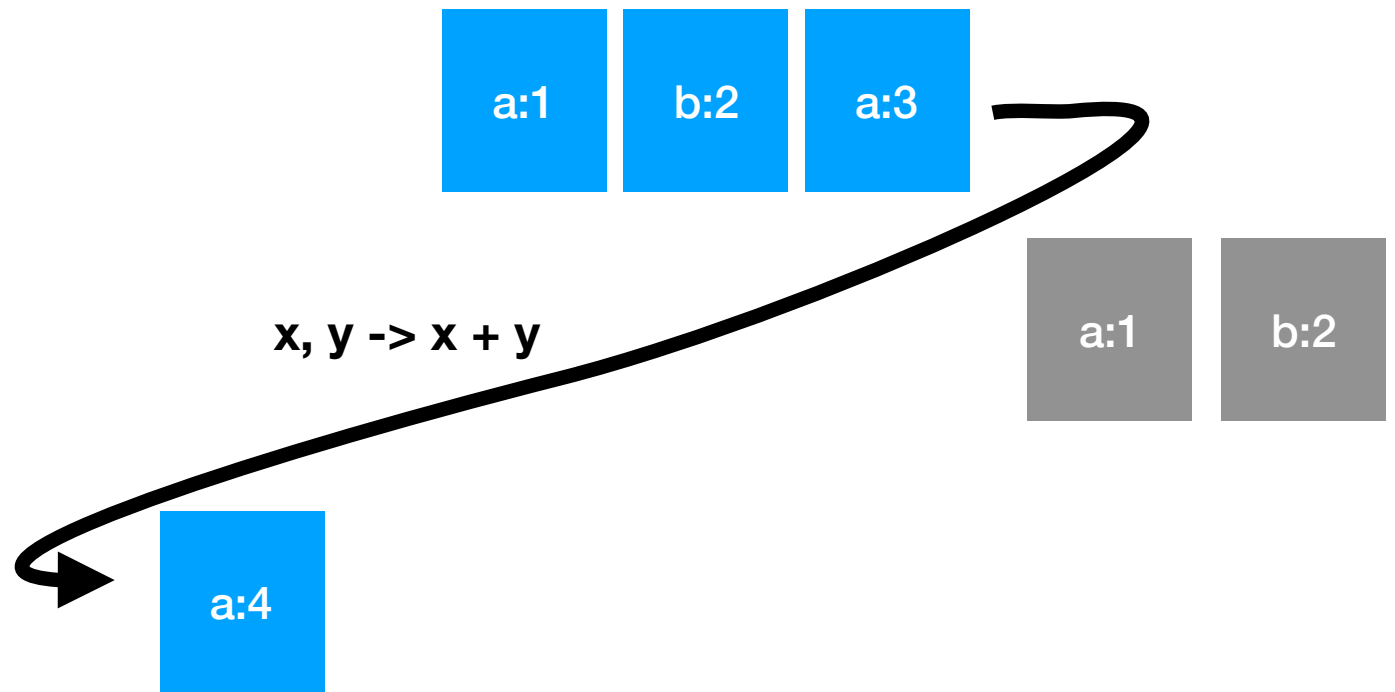
# Fundamental stream query operators

1. Filter
2. Map
- 3. Join**
4. Aggregate
5. Zip



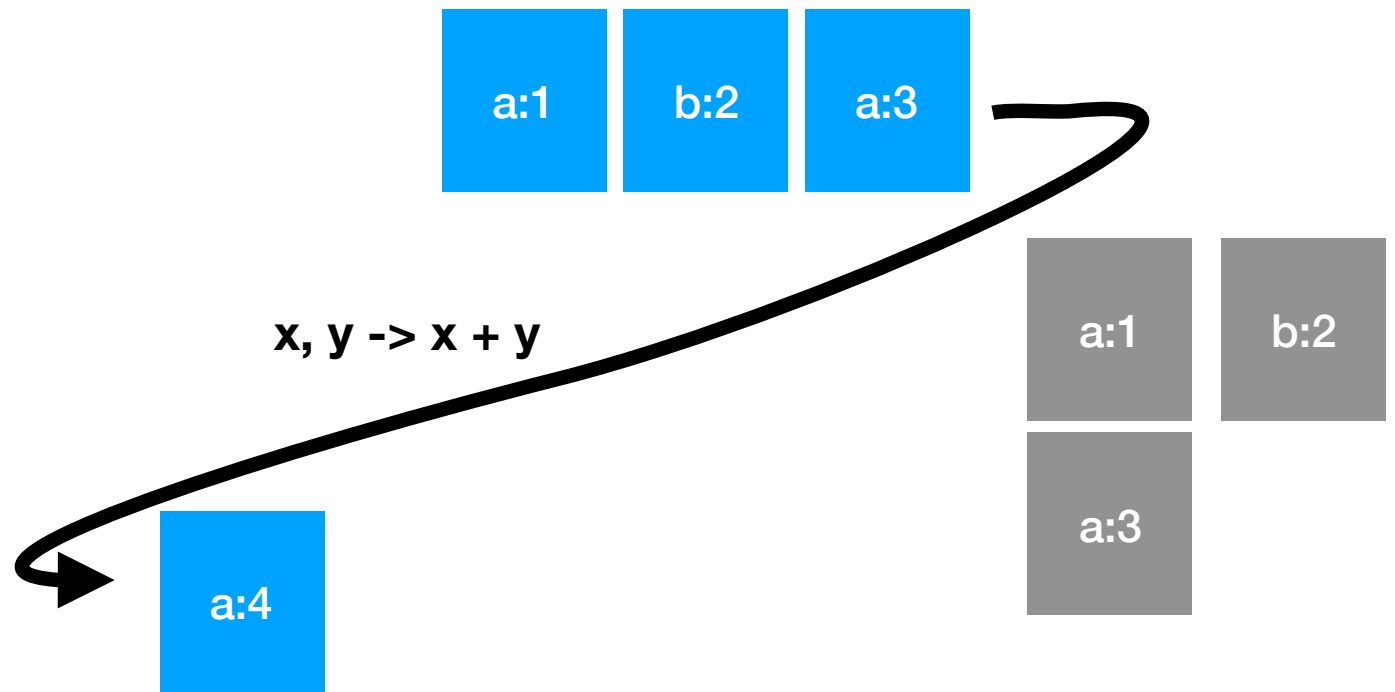
# Fundamental stream query operators

1. Filter
2. Map
- 3. Join**
4. Aggregate
5. Zip



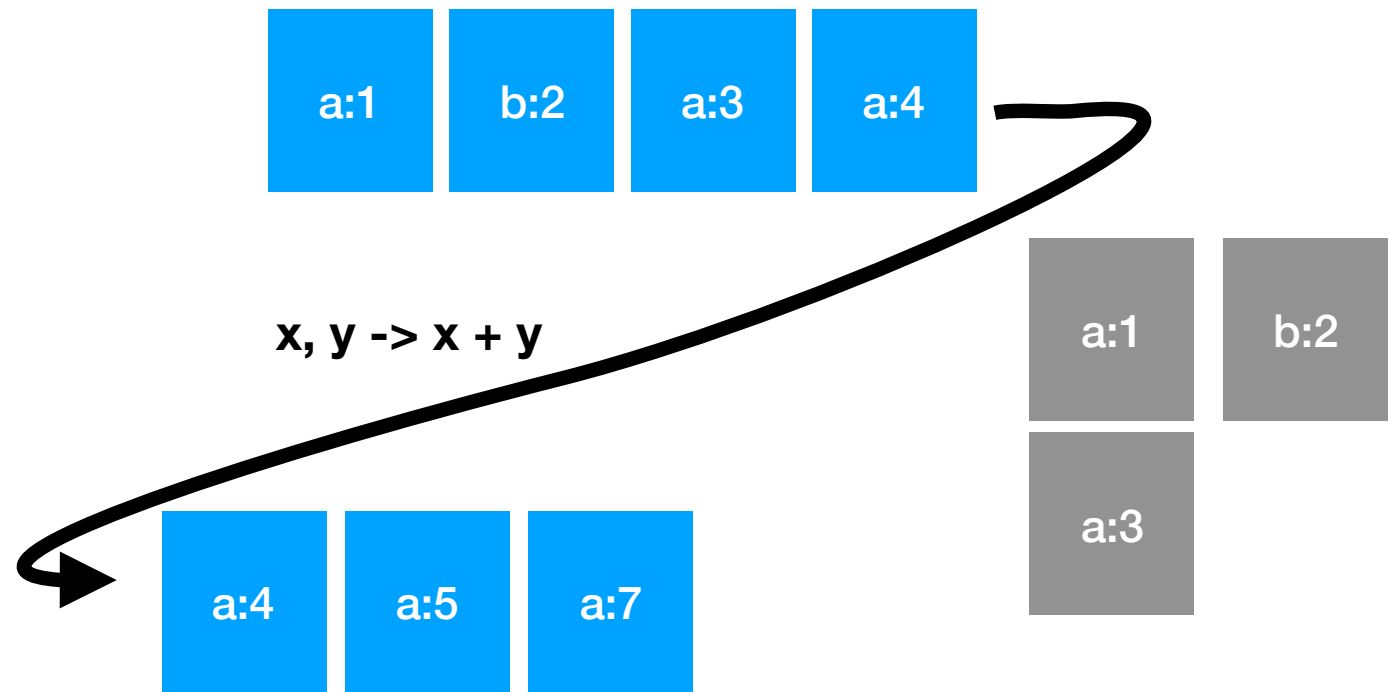
# Fundamental stream query operators

1. Filter
2. Map
- 3. Join**
4. Aggregate
5. Zip



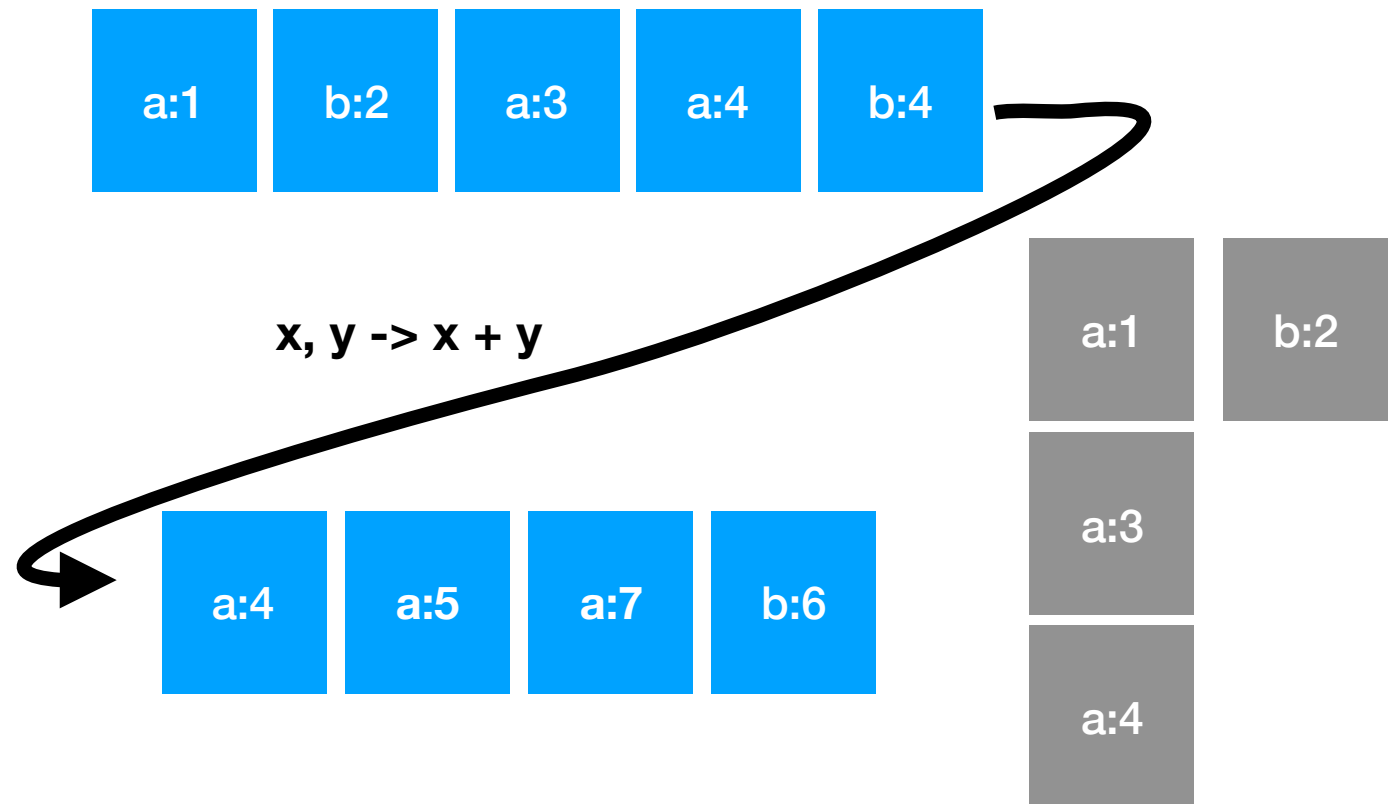
# Fundamental stream query operators

1. Filter
2. Map
- 3. Join**
4. Aggregate
5. Zip



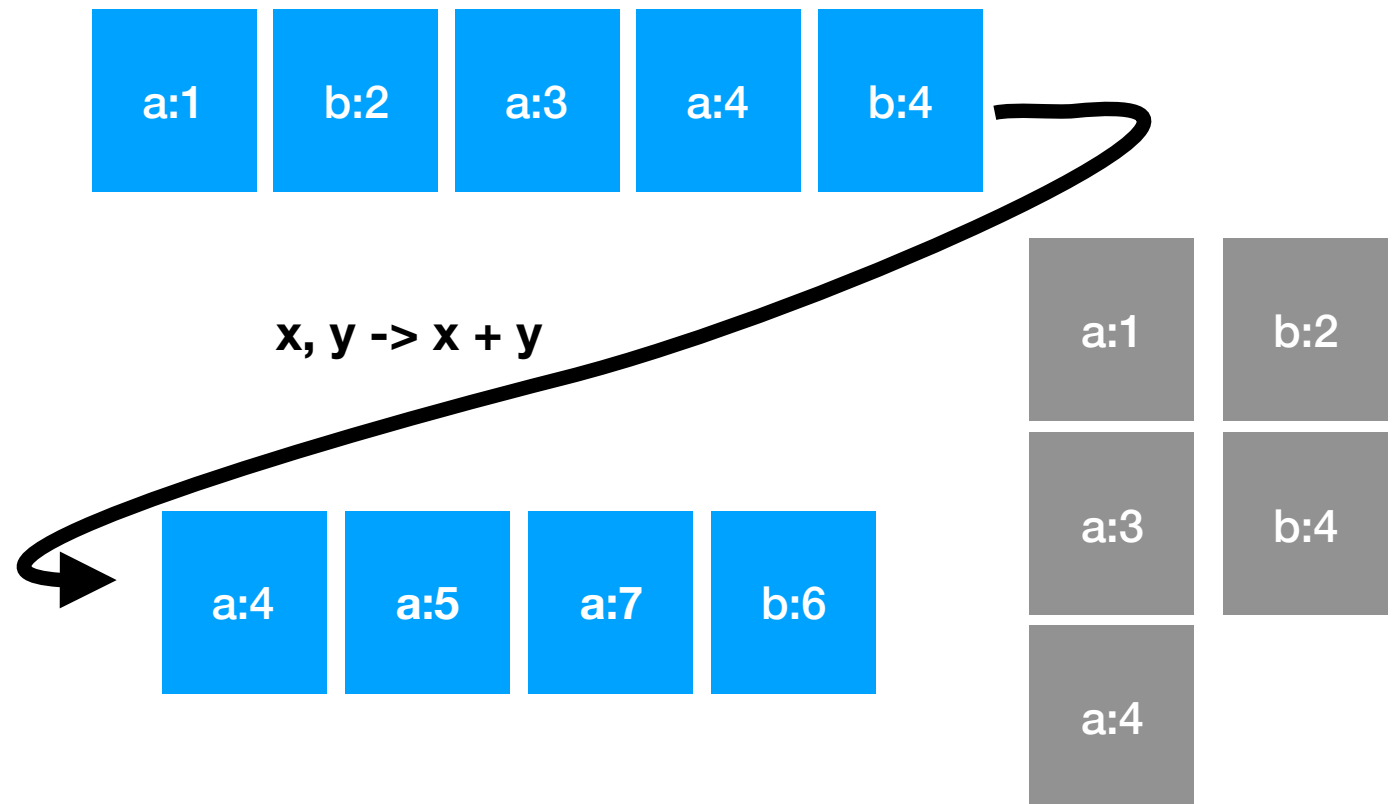
# Fundamental stream query operators

1. Filter
2. Map
- 3. Join**
4. Aggregate
5. Zip



# Fundamental stream query operators

1. Filter
2. Map
- 3. Join**
4. Aggregate
5. Zip



# Fundamental stream query operators

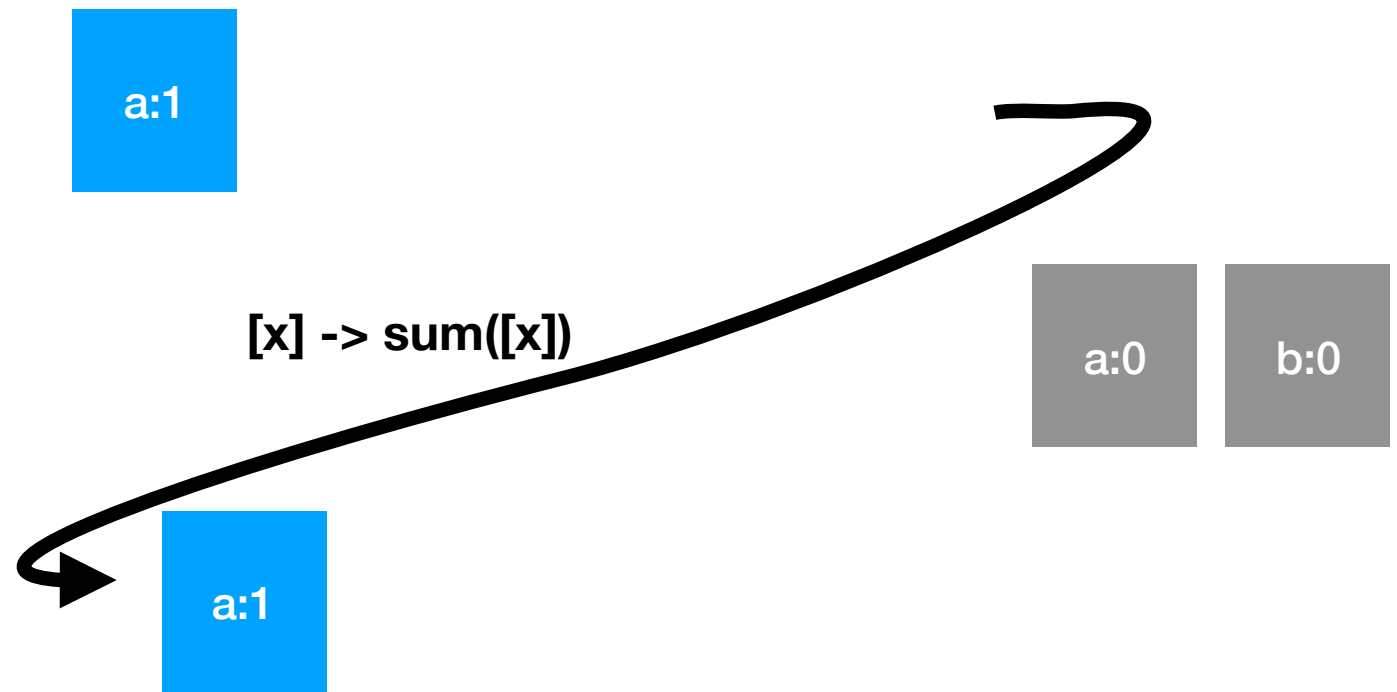
1. Filter

2. Map

3. Join

**4. Aggregate**

5. Zip





# Fundamental stream query operators

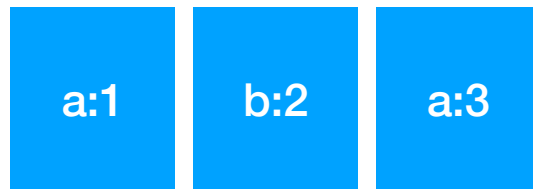
1. Filter

2. Map

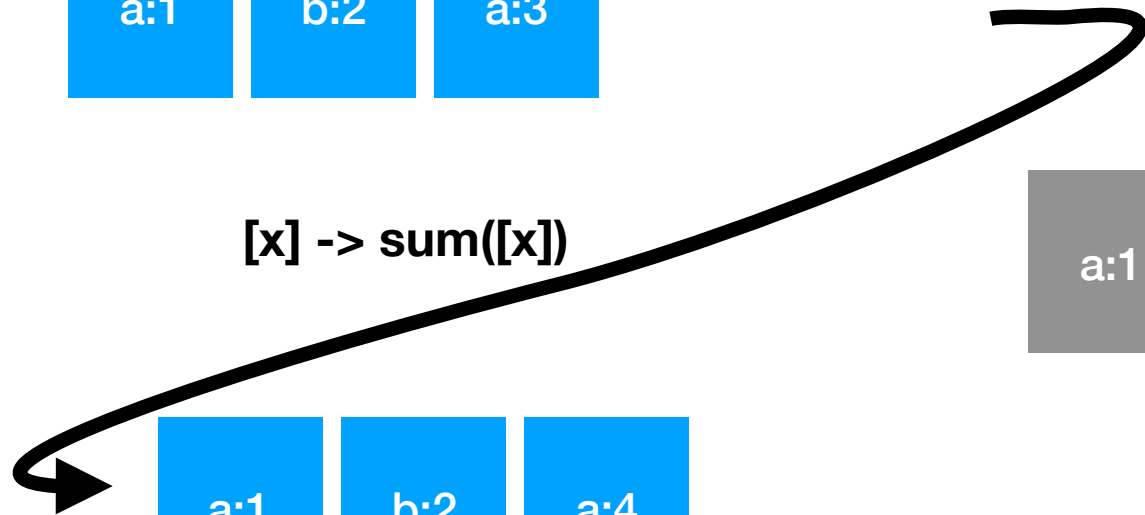
3. Join

**4. Aggregate**

5. Zip



$[x] \rightarrow \text{sum}([x])$



# Fundamental stream query operators

1. Filter

2. Map

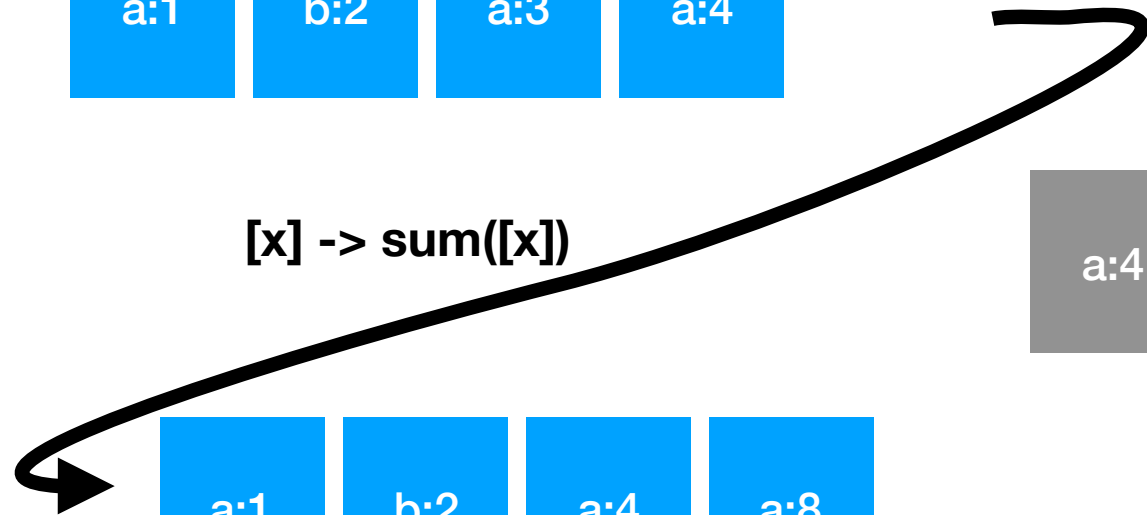
3. Join

**4. Aggregate**

5. Zip

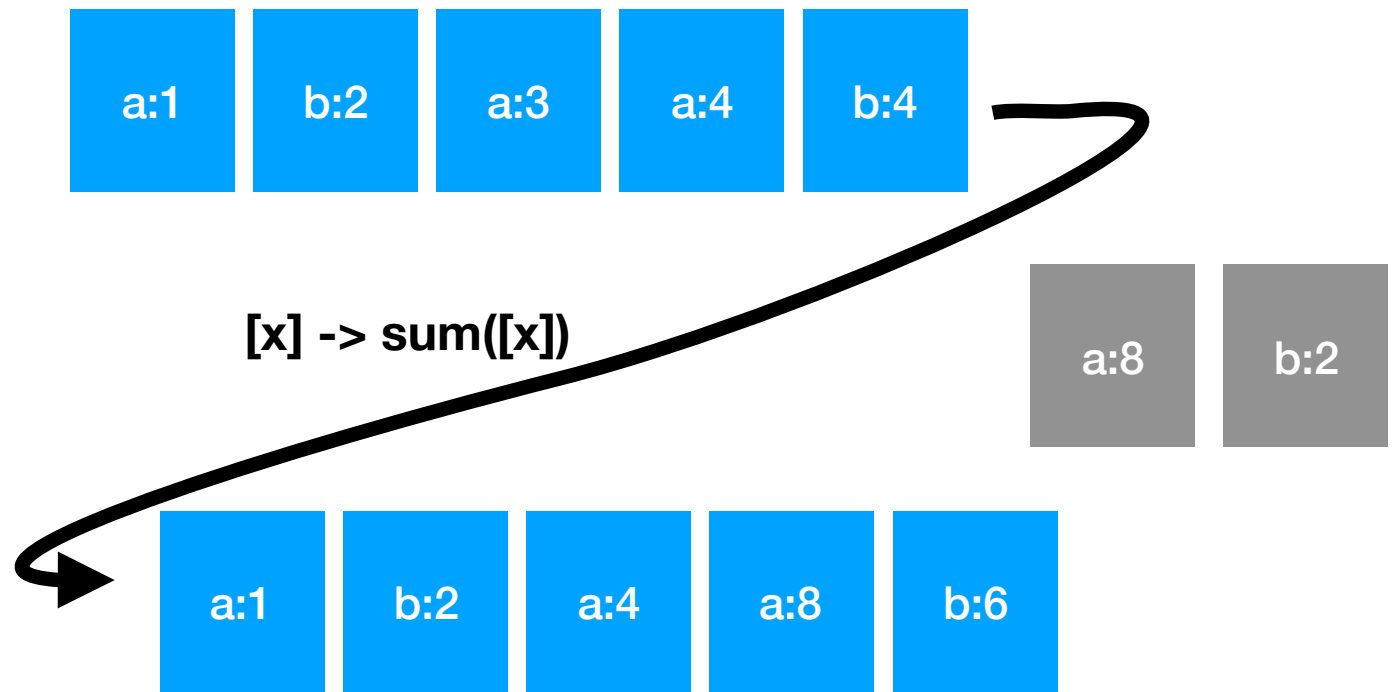


$[x] \rightarrow \text{sum}([x])$



# Fundamental stream query operators

1. Filter
2. Map
3. Join
- 4. Aggregate**
5. Zip



# Fundamental stream query operators

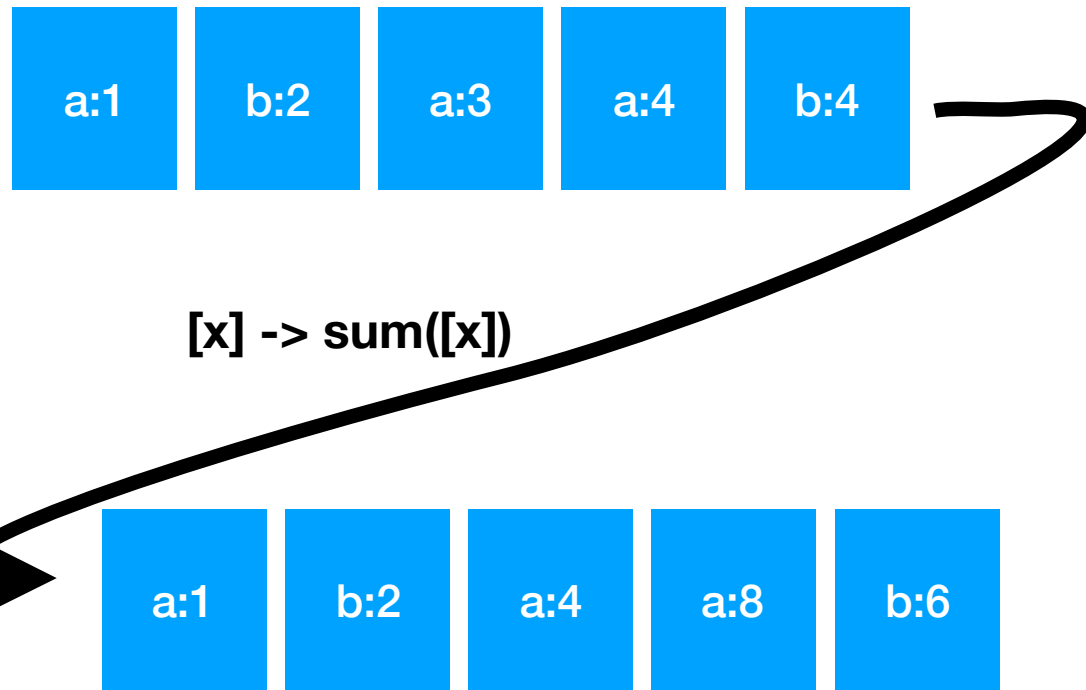
1. Filter

2. Map

3. Join

**4. Aggregate**

5. Zip



# Fundamental stream query operators

1. Filter
2. Map
3. Join
4. Aggregate
5. **Zip**

b:2	b:1	a:3	b:1	a:1
a:1	b:2	a:3	a:4	b:4

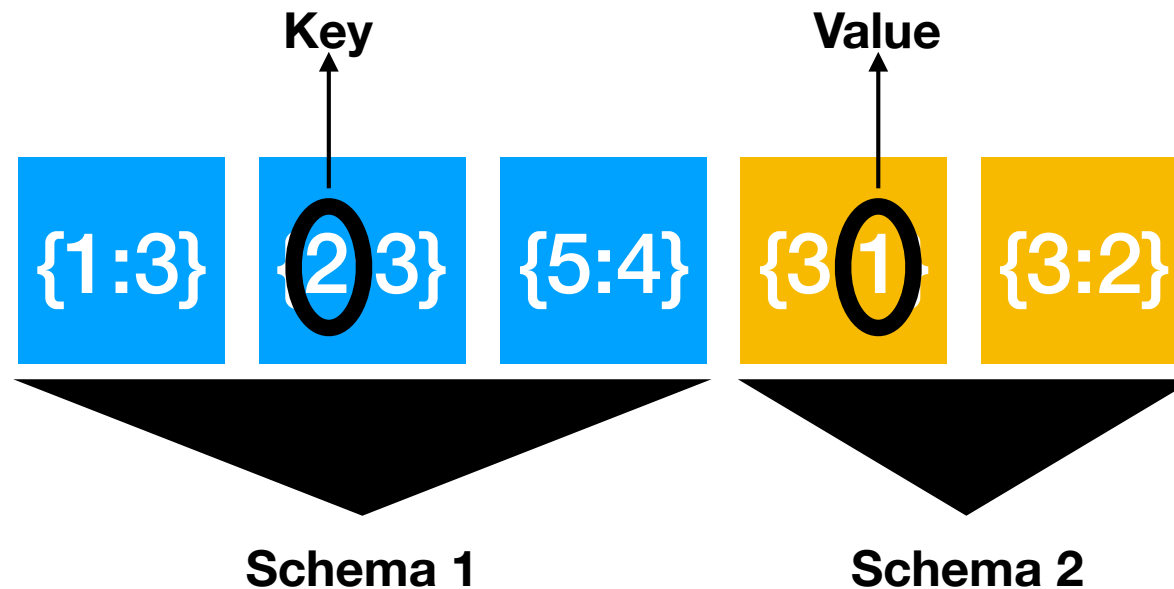
$(k1, v1), (k2, v2) \rightarrow (k1 + k2, v1 + v2)$

c:3	b:3	b:6	c:5	c:5
-----	-----	-----	-----	-----

# Fundamental stream query operators

1. Filter
- 2. Map**
- 3. Join**
- 4. Aggregate**
5. Zip

# P4 Implementation: Stream Objects

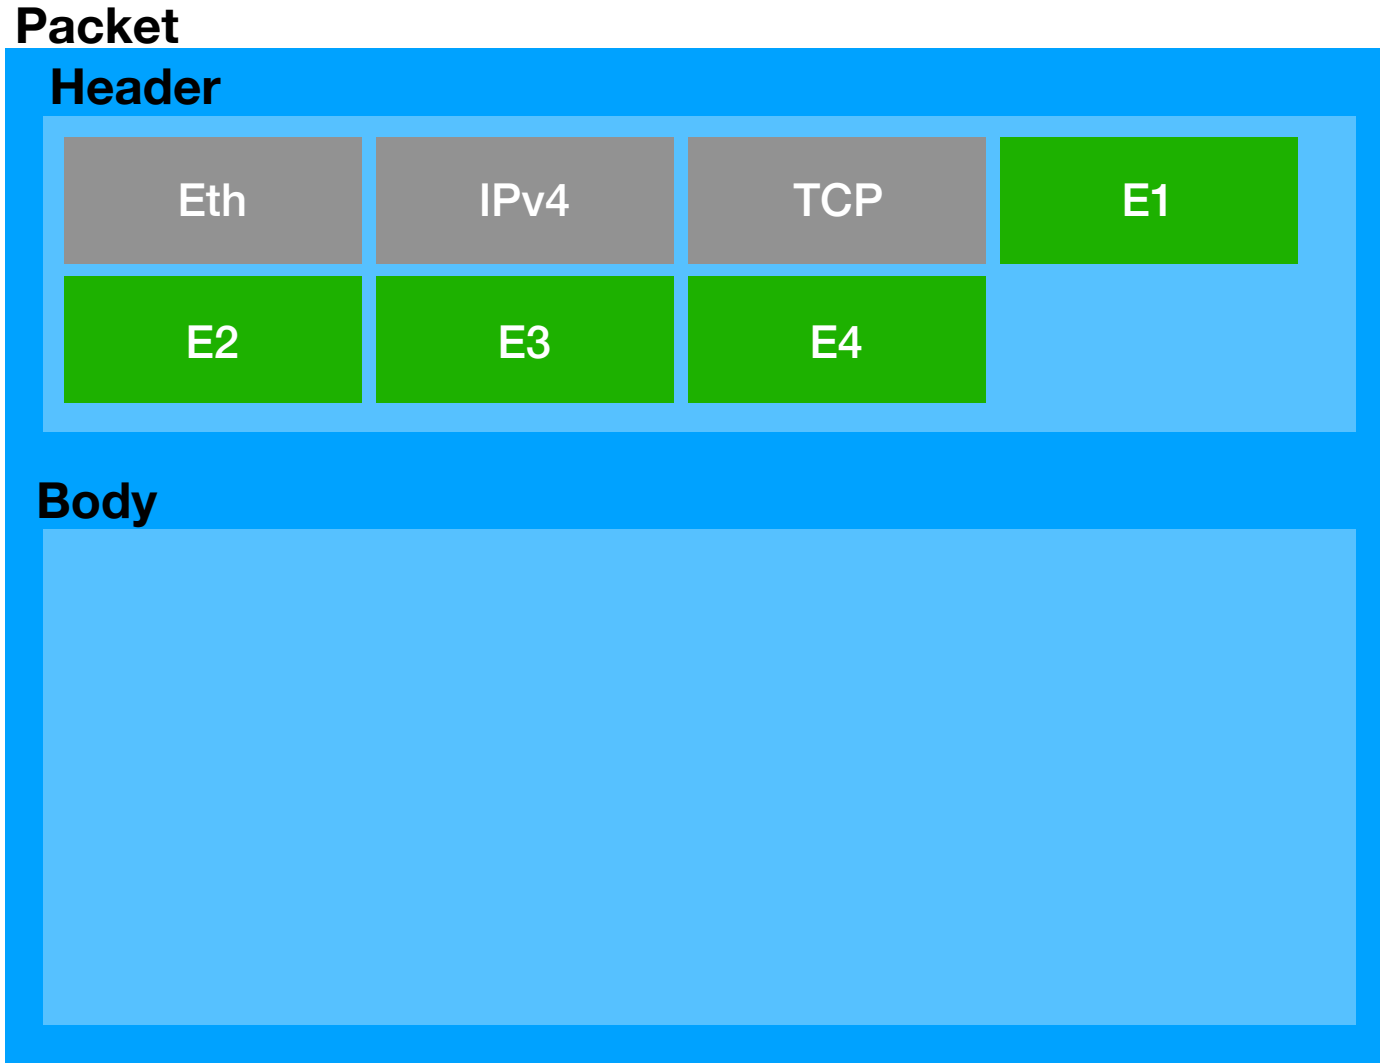


# P4 Implementation: Header Format

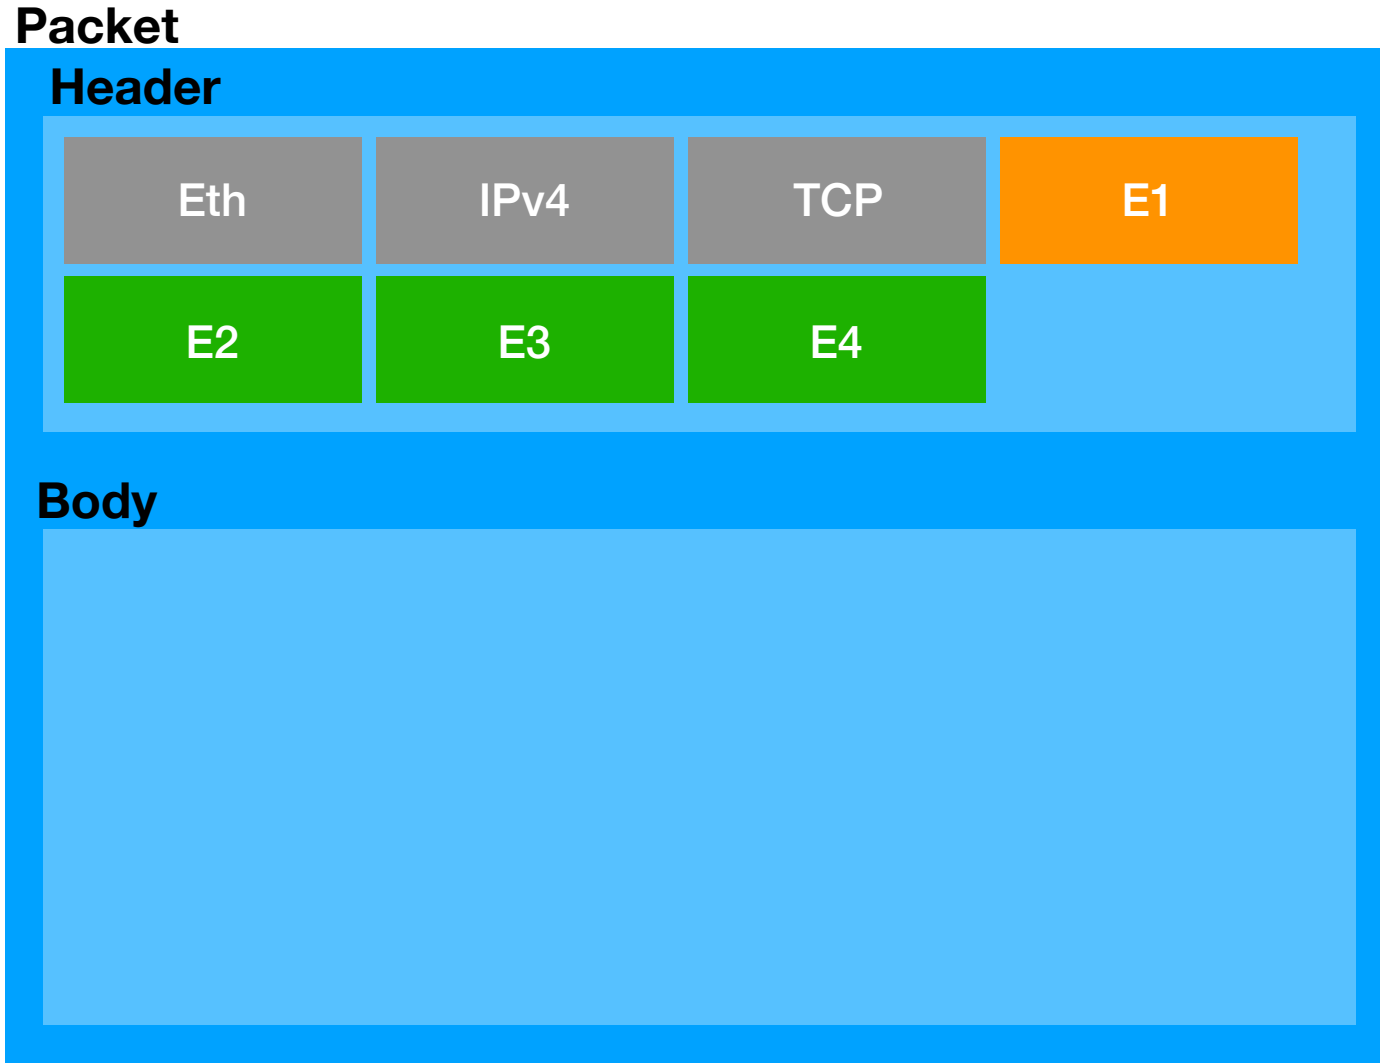
```
header entry_t {  
    bit<32> schema;  
    bit<32> key;  
    bit<32> val;  
    bit<8> unprocessed;  
}
```



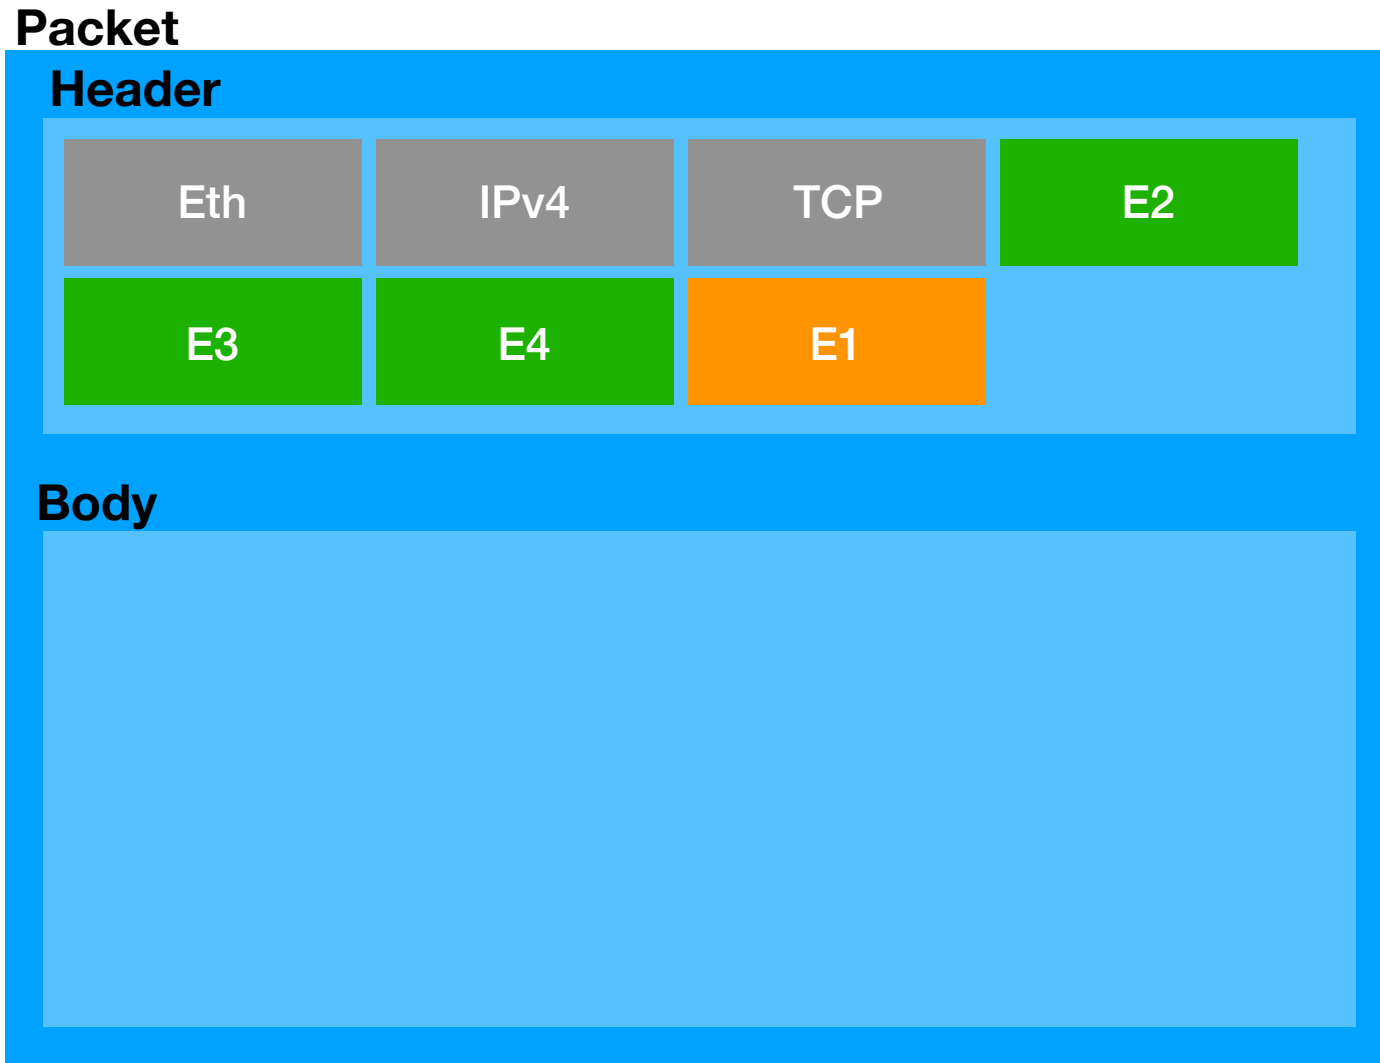
# Always process first stream object on stack



# Always process first stream object on stack



# Circulate stack to get next element



**recirculate** packets and use  
temporary header field to rotate stack

```
struct headers {  
    ethernet_t    ethernet;  
    ipv4_t        ipv4;  
    tcp_t         tcp;  
    entry_count_t entry_count;  
    entry_t[10] entry;  
    entry_t entry_swap;  
}
```

# P4 Implementation: Representing Stream Operations

```
{
  "table": "MyIngress.stream_ops",
  "match": {
    "hdr.entry[0].schema": 0,
    "meta.lineNo": 0
  },
  "action_name": "MyIngress.add",
  "action_params": {
    "i": 1
  }
}
```

# P4 Implementation: Representing Stream Operations

```
{
  "table": "MyIngress.stream_ops",
  "match": {
    "hdr.entry[0].schema": 0,
    "meta.lineNo": 0
  },
  "action_name": "MyIngress.add",
  "action_params": {
    "i": 1
  }
}
```

# P4 Implementation: Representing Stream Operations

```
{  
  "table": "MyIngress.stream_ops",  
  "match": {  
    "hdr.entry[0].schema": 0,  
    "meta.lineNo": 0  
  },  
  "action_name": "MyIngress.add",  
  "action_params": {  
    "i": 1  
  }  
}
```

# P4 Implementation: Representing Stream Operations

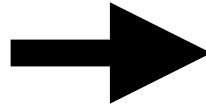
```
{
  "table": "MyIngress.stream_ops",
  "match": {
    "hdr.entry[0].schema": 0,
    "meta.lineNo": 0
  },
  "action_name": "MyIngress.add",
  "action_params": {
    "i": 1
  }
},
{
  "table": "MyIngress.stream_ops",
  "match": {
    "hdr.entry[0].schema": 0,
    "meta.lineNo": 1
  },
  "action_name": "MyIngress.add",
  "action_params": {
    "i": 1
  }
}
```



# Map operator in action

```
###[ TCP ]###
sport      = 56169
dport      = 4660
seq        = 0
ack        = 0
dataofs    = 5L
reserved   = 0L
flags      = S
window     = 8192
chksum     = 0x8a26
urgptr     = 0
options    = []
###[ KeyValCount ]###
count      = 1
###[ KeyValPair ]###
schema     = 0
key        = 3
val        = 7
unprocessed= 1L
```

**{ 3 : 7 }**



```
###[ TCP ]###
sport      = 56169
dport      = 4660
seq        = 0
ack        = 0
dataofs    = 5L
reserved   = 0L
flags      = S
window     = 8192
chksum     = 0x8a26
urgptr     = 0
options    = []
###[ KeyValCount ]###
count      = 1
###[ KeyValPair ]###
schema     = 0
key        = 3
val        = 9
unprocessed= 0L
```

**{ 3 : 9 }**

# P4 Implementation: Representing Stream Operations

```
{
  "table": "MyIngress.stream_ops",
  "match": {
    "hdr.entry[0].schema": 1,
    "meta.lineNo": 0
  },
  "action_name":
  "MyIngress.key_window_aggregate",
  "action_params": { }
}
```

# Aggregate operator in action

```
###[ TCP ]###
sport      = 56169
dport      = 4660
seq        = 0
ack        = 0
dataofs    = 5L
reserved   = 0L
flags      = S
window     = 8192
chksum     = 0x8a26
urgptr     = 0
options    = []
###[ KeyValCount ]###
count      = 1
###[ KeyValPair ]###
schema     = 0
key        = 3
val        = 7
unprocessed= 1L
```

**{ 4 : 7 }**

```
###[ TCP ]###
sport      = 52165
dport      = 4660
seq        = 0
ack        = 0
dataofs    = 5L
reserved   = 0L
flags      = S
window     = 8192
chksum     = 0x99c7
urgptr     = 0
options    = []
###[ KeyValCount ]###
count      = 1
###[ KeyValPair ]###
schema     = 1
key        = 4
val        = 8
unprocessed= 1L
```

**{ 4 : 8 }**

```
###[ TCP ]###
sport      = 62853
dport      = 4660
seq        = 0
ack        = 0
dataofs    = 5L
reserved   = 0L
flags      = S
window     = 8192
chksum     = 0x7006
urgptr     = 0
options    = []
###[ KeyValCount ]###
count      = 1
###[ KeyValPair ]###
schema     = 1
key        = 4
val        = 9
unprocessed= 1L
```

**{ 4 : 9 }**



```
###[ TCP ]###
sport      = 62853
dport      = 4660
seq        = 0
ack        = 0
dataofs    = 5L
reserved   = 0L
flags      = S
window     = 8192
chksum     = 0x7006
urgptr     = 0
options    = []
###[ KeyValCount ]###
count      = 1
###[ KeyValPair ]###
schema     = 1
key        = 4
val        = 24
unprocessed= 0L
```

**{ 4 : 24 }**

# P4 Implementation: Representing Stream Operations

```
{
  "table": "MyIngress.stream_ops",
  "match": {
    "hdr.entry[0].schema": 2,
    "meta.lineNo": 0
  },
  "action_name":
  "MyIngress.join_sum",
  "action_params": { }
}
```

# Join operator in action

```
###[ KeyValCount ]###  
count = 1  
###[ KeyValPair ]###  
schema = 2  
key = 1  
val = 1  
unprocessed= 1L
```

{1:1}

```
###[ KeyValCount ]###  
count = 1  
###[ KeyValPair ]###  
schema = 2  
key = 2  
val = 2  
unprocessed= 1L
```

{2:2}

```
###[ KeyValCount ]###  
count = 1  
###[ KeyValPair ]###  
schema = 2  
key = 1  
val = 3  
unprocessed= 1L
```

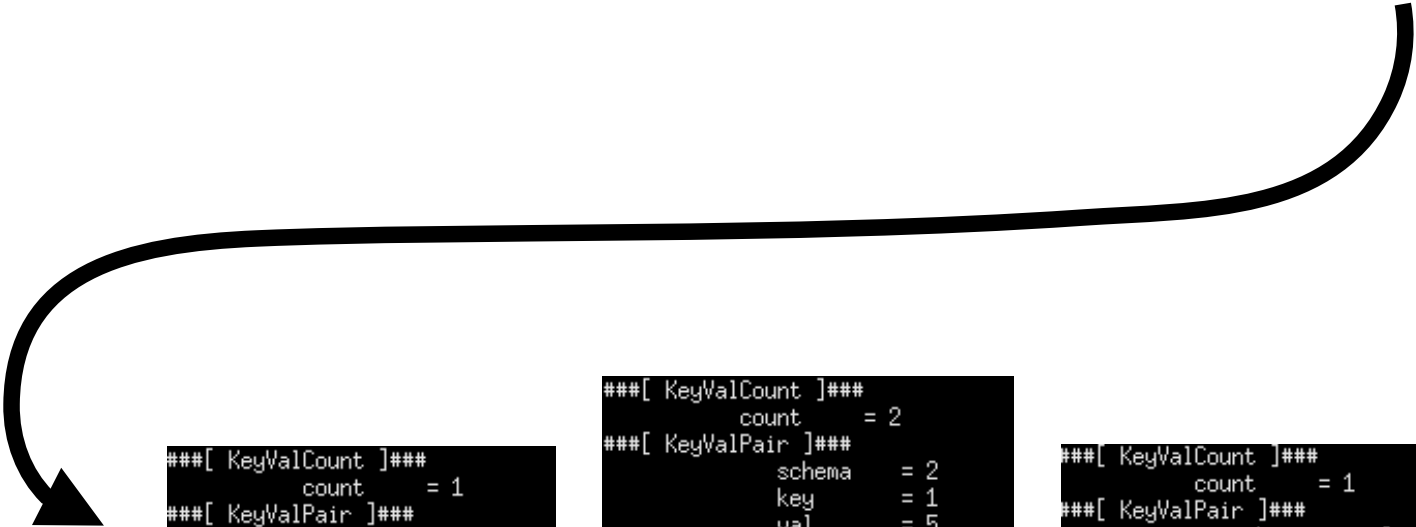
{1:3}

```
###[ KeyValCount ]###  
count = 1  
###[ KeyValPair ]###  
schema = 2  
key = 1  
val = 4  
unprocessed= 1L
```

{1:4}

```
###[ KeyValCount ]###  
count = 1  
###[ KeyValPair ]###  
schema = 2  
key = 2  
val = 4  
unprocessed= 1L
```

{2:4}



```
###[ KeyValCount ]###  
count = 1  
###[ KeyValPair ]###  
schema = 2  
key = 1  
val = 4  
unprocessed= 0L
```

{1:4}

```
###[ KeyValCount ]###  
count = 2  
###[ KeyValPair ]###  
schema = 2  
key = 1  
val = 5  
unprocessed= 0L  
###[ KeyValPair ]###  
schema = 2  
key = 1  
val = 7  
unprocessed= 0L
```

{1:5}, {1:7}

```
###[ KeyValCount ]###  
count = 1  
###[ KeyValPair ]###  
schema = 2  
key = 2  
val = 6  
unprocessed= 0L
```

{2:6}

# Next steps

- Implement Stream QL compiler
- Compare to python stream processing implementation